

AI Development Notes

Ollama is a tool used to run Large Language Models locally on a computer.

Qwen is a language model that can answer questions and generate text.

LLaVA is a vision language model that can analyze images and describe their contents.

Whisper is a speech-to-text model used to convert audio recordings into text.

Retrieval Augmented Generation (RAG) combines document retrieval with language models to provide more accurate answers.

Vector databases such as ChromaDB store document embeddings and help retrieve relevant information.

Project Goal:

Build a local multimodal assistant that can analyze images, transcribe audio, search PDFs, and generate summaries.